

Neural Network Laboratory Record

B.E. (AI & DS) – VI Semester

Name: Rahul Raj

Roll No: 24UADS1041

Experiment Title: Implementation of Multi-Layer Perceptron for XOR Problem

1. Objective

To implement a multi-layer perceptron (MLP) network with one hidden layer using NumPy in Python. Demonstrate that it can learn the XOR Boolean function.

2. Introduction

A Multi-Layer Perceptron (MLP) is a feed-forward artificial neural network consisting of an input layer, one or more hidden layers, and an output layer. Unlike a single-layer perceptron, an MLP can solve non-linear problems using hidden neurons and non-linear activation functions such as sigmoid.

The XOR function is a classic example of a non-linearly separable problem. A single perceptron cannot solve it, but an MLP can learn it by forming a non-linear decision boundary.

3. Dataset Used

| Input 1 | Input 2 | Output |
|---------|---------|--------|
| 0 | 0 | 0 |
| 0 | 1 | 1 |
| 1 | 0 | 1 |
| 1 | 1 | 0 |

4. Python Implementation

```
# -----
```

```
# MLP for XOR Problem using NumPy
```

```
# -----
```

```
import numpy as np
```

```
import matplotlib.pyplot as plt
```

```
# -----
```

```
# 1 XOR Dataset
```

```
# -----
```

```
X = np.array([
```

```
    [0, 0],
```

```
    [0, 1],
```

```
    [1, 0],
```

```
    [1, 1]
```

```
])
```

```
y = np.array([[0], [1], [1], [0]])
```

```
# -----
```

```
# 2 Activation Function
```

```
# -----
```

```
def sigmoid(x):
```

```
    return 1 / (1 + np.exp(-x))
```

```
def sigmoid_derivative(x):
```

```
    return x * (1 - x)
```

```
# -----
```

```
# 3 Initialize Network Parameters
```

```
# -----
```

```
np.random.seed(42)
```

```
input_size = 2
```

```
hidden_size = 4
```

```
output_size = 1
```

```
# Weights  
  
W1 = np.random.randn(input_size, hidden_size)  
  
b1 = np.zeros((1, hidden_size))
```

```
W2 = np.random.randn(hidden_size, output_size)  
  
b2 = np.zeros((1, output_size))
```

```
# -----
```

```
# 4 Hyperparameters
```

```
# -----
```

```
learning_rate = 0.1
```

```
epochs = 10000
```

```
losses = []
```

```
# 5 Training MLP (Backpropagation)
```

```
# -----
```

```
for epoch in range(epochs):
```

```
    # Forward Pass
```

```
    hidden_input = np.dot(X, W1) + b1
```

```
    hidden_output = sigmoid(hidden_input)
```

```
    final_input = np.dot(hidden_output, W2) + b2
```

```
    output = sigmoid(final_input)
```

```
# Compute Loss (MSE)
```

```
    loss = np.mean((y - output)**2)
```

```
    losses.append(loss)
```

```
# Backpropagation
```

```
d_output = (y - output) * sigmoid_derivative(output)

d_hidden = d_output.dot(W2.T) * sigmoid_derivative(hidden_output)

# Update Weights and Biases

W2 += hidden_output.T.dot(d_output) * learning_rate

b2 += np.sum(d_output, axis=0, keepdims=True) * learning_rate

W1 += X.T.dot(d_hidden) * learning_rate

b1 += np.sum(d_hidden, axis=0, keepdims=True) * learning_rate

# Print loss every 1000 epochs

if epoch % 1000 == 0:

    print(f"Epoch {epoch}, Loss: {loss:.6f}")

# -----

# 6 Final Predictions

# -----

print("\nFinal Predictions (after training):")

print(output)

# -----

# 7 Plot Loss Curve

# -----

plt.plot(losses)

plt.title("Loss vs Epochs")

plt.xlabel("Epochs")

plt.ylabel("Mean Squared Error")

plt.show()

# -----
```

```

# 8 Plot Decision Boundary

# ----

def plot_decision_boundary():

    x_min, x_max = -0.5, 1.5
    y_min, y_max = -0.5, 1.5
    h = 0.01

    xx, yy = np.meshgrid(np.arange(x_min, x_max, h),
                         np.arange(y_min, y_max, h))

    grid = np.c_[xx.ravel(), yy.ravel()]

    hidden = sigmoid(np.dot(grid, W1) + b1)
    out = sigmoid(np.dot(hidden, W2) + b2)

    Z = out.reshape(xx.shape)

    plt.contourf(xx, yy, Z, levels=[0, 0.5, 1], alpha=0.3, colors=['#FFAAAA','#AAAAFF'])

    plt.scatter(X[:,0], X[:,1], c=y.flatten(), edgecolors='k', s=100)

    plt.title("Decision Boundary")
    plt.xlabel("Input 1")
    plt.ylabel("Input 2")
    plt.show()

plot_decision_boundary()

```

5. Output

Epoch 0, Loss: 0.283190
 Epoch 1000, Loss: 0.245226
 Epoch 2000, Loss: 0.212412
 Epoch 3000, Loss: 0.150331
 Epoch 4000, Loss: 0.057156
 Epoch 5000, Loss: 0.020929
 Epoch 6000, Loss: 0.010685

Epoch 7000, Loss: 0.006679

Epoch 8000, Loss: 0.004683

Epoch 9000, Loss: 0.003527

Final Predictions (after training):

[[0.03730284]

[0.9491398]

[0.94480964]

[0.06425255]]

6. Conclusion

The Multi-Layer Perceptron successfully learned the XOR Boolean function. The decreasing loss over epochs indicates stable convergence. This experiment demonstrates that hidden layers enable neural networks to solve non-linear problems.

7. Graphical Results

Figure 1: Loss Curve




Figure 2: Decision Boundary

